

Part 3: Deadlocks

Overview


- ✦ Resources
- ✦ Why do deadlocks occur?
- ✦ Dealing with deadlocks
 - Ignoring them: ostrich algorithm
 - Detecting & recovering from deadlock
 - Avoiding deadlock
 - Preventing deadlock

Resources

- ✦ Resource: something a process uses
 - Usually limited (at least somewhat)
- ✦ Examples of computer resources
 - Printers
 - Semaphores / locks
 - Tables (in a database)
- ✦ Processes need access to resources in reasonable order
- ✦ Two types of resources:
 - Preemptable resources: can be taken away from a process with no ill effects
 - Nonpreemptable resources: will cause the process to fail if taken away

When do deadlocks happen?

- ✦ Suppose
 - Process 1 holds resource A and requests resource B
 - Process 2 holds B and requests A
 - Both can be blocked, with neither able to proceed
- ✦ Deadlocks occur when ...
 - Processes are granted exclusive access to devices or software constructs (resources)
 - Each deadlocked process needs a resource held by another deadlocked process



Using resources

- ✦ Sequence of events required to use a resource
 - Request the resource
 - Use the resource
 - Release the resource
- ✦ Can't use the resource if request is denied
 - Requesting process has options
 - Block and wait for resource
 - Continue (if possible) without it: may be able to use an alternate resource
 - Process fails with error code
 - Some of these may be able to prevent deadlock...

What is a deadlock?

- ◆ Formal definition:

“A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.”

- ◆ Usually, the event is release of a currently held resource


- ◆ In deadlock, none of the processes can

- Run
- Release resources
- Be awakened

Four conditions for deadlock

- ✦ Mutual exclusion
 - Each resource is assigned to at most one process
- ✦ Hold and wait
 - A process holding resources can request more resources
- ✦ No preemption
 - Previously granted resources cannot be forcibly taken away
- ✦ Circular wait
 - There must be a circular chain of 2 or more processes where each is waiting for a resource held by the next member of the chain

Resource allocation graphs



- ✦ Resource allocation modeled by directed graphs
- ✦ Example 1:
 - Resource R assigned to process A
- ✦ Example 2:
 - Process B is requesting / waiting for resource S
- ✦ Example 3:
 - Process C holds T, waiting for U
 - Process D holds U, waiting for T
 - C and D are in deadlock!



Dealing with deadlock

- ✦ How can the OS deal with deadlock?
 - Ignore the problem altogether!
 - Hopefully, it'll never happen...
 - Detect deadlock & recover from it
 - Dynamically avoid deadlock
 - Careful resource allocation
 - Prevent deadlock
 - Remove at least one of the four necessary conditions
- ✦ We'll explore these tradeoffs

Getting into deadlock



A

Acquire R
Acquire S
Release R
Release S





B

Acquire S
Acquire T
Release S
Release T



C

Acquire T
Acquire R
Release T
Release R



Not getting into deadlock...

- ✦ Many situations may result in deadlock (but don't have to)
 - In previous example, A could release R before C requests R, resulting in no deadlock
 - Can we always get out of it this way?
- ✦ Find ways to:
 - Detect deadlock and reverse it
 - Stop it from happening in the first place


The Ostrich Algorithm

- ✦ Pretend there's no problem
- ✦ Reasonable if
 - Deadlocks occur very rarely
 - Cost of prevention is high
- ✦ UNIX and Windows take this approach
 - Resources (memory, CPU, disk space) are plentiful
 - Deadlocks over such resources rarely occur
 - Deadlocks typically handled by rebooting
- ✦ Trade off between convenience and correctness

Detecting deadlocks using graphs

- ◆ Process holdings and requests in the table and in the graph (they're equivalent)
- ◆ Graph contains a cycle \Rightarrow deadlock!
 - Easy to pick out by looking at it (in this case)
 - Need to mechanically detect deadlock
- ◆ Not all processes are deadlocked (A, C, F not in deadlock)

Process	Holds	Wants
A	R	S
B		T
C		S
D	U	S,T
E	T	V
F	W	S
G	V	U



Deadlock detection algorithm

- ✦ General idea: try to find cycles in the resource allocation graph
- ✦ Algorithm: depth-first search at each node
 - Mark arcs as they're traversed
 - Build list of visited nodes
 - If node to be added is already on the list, a cycle exists!
- ✦ Cycle \Rightarrow deadlock

```
For each node N in the graph {  
    Set L = empty list  
    unmark all arcs  
    Traverse (N,L)  
}  
If no deadlock reported by now,  
there isn't any  
  
define Traverse (C,L) {  
    If C in L, report deadlock!  
    Add C to L  
    For each unmarked arc from C {  
        Mark the arc  
        Set A = arc destination  
        /* NOTE: L is a  
           local variable */  
        Traverse (A,L)  
    }  
}
```

Resources with multiple instances

- ✦ Previous algorithm only works if there's one instance of each resource
- ✦ If there are multiple instances of each resource, we need a different method
 - Track current usage and requests for each process
 - To detect deadlock, try to find a scenario where all processes can finish
 - If no such scenario exists, we have deadlock

Deadlock detection algorithm

	A	B	C	D
Avail	2	3	0	1

Hold

Process	A	B	C	D
1	0	3	0	0
2	1	0	1	1
3	0	2	1	0
4	2	2	3	0

Want

Process	A	B	C	D
1	3	2	1	0
2	2	2	0	0
3	3	5	3	1
4	0	4	1	1


```
current=avail;
for (j = 0; j < N; j++) {
    for (k=0; k<N; k++) {
        if (finished[k])
            continue;
        if (want[k] < current) {
            finished[k] = 1;
            current += hold[k];
            break;
        }
        if (k==N) {
            printf "Deadlock!\n";
            // finished[k]==0 means
            process
            // is in the deadlock
            break;
        }
    }
}
```

Note: want[j],hold[j],current,avail are arrays!

Recovering from deadlock

- ✦ Recovery through preemption
 - Take a resource from some other process
 - Depends on nature of the resource and the process
- ✦ Recovery through rollback
 - Checkpoint a process periodically
 - Use saved state to restart the process if it's in deadlock
 - May present a problem if the process affects lots of “external” things
- ✦ Recovery through killing processes
 - Crudest but simplest way to break a deadlock: kill one of the processes in the deadlock cycle
 - Other processes can get its resources
 - Try to choose a process that can be rerun from the start
 - Pick one that hasn't run too far already

Resource trajectories



Two process resource trajectories

Safe and unsafe states

	Has	Max		Has	Max		Has	Max		Has	Max		Has	Max		Has	Max
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	-	B	0	-	B	0	-	B	0	-
C	2	7	C	2	7	C	2	7	C	7	7	C	0	-	C	0	-
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7					

Demonstration that the first state is safe

	Has	Max		Has	Max		Has	Max		Has	Max
A	3	9	A	4	9	A	4	9	A	4	9
B	2	4	B	2	4	B	4	4	B	0	-
C	2	7	C	2	7	C	2	7	C	2	7
Free: 3			Free: 2			Free: 0			Free: 4		

Demonstration that the second state is unsafe

Banker's Algorithm for a single resource

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

Any sequence finishes

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

C,B,A,D finishes

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

Deadlock (unsafe state)

- ✦ Bankers' algorithm: before granting a request, ensure that a sequence exists that will allow all processes to complete
 - Use previous methods to find such a sequence
 - If a sequence exists, allow the requests
 - If there's no such sequence, deny the request
- ✦ Can be slow: must be done on each request!

Banker's Algorithm for multiple resources

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Scanners	CD ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

E = (6342)
P = (5322)
A = (1020)

Example of banker's algorithm with multiple resources

Preventing deadlock

- ✦ Deadlock can be completely prevented!
- ✦ Ensure that at least one of the conditions for deadlock never occurs
 - Mutual exclusion
 - Circular wait
 - Hold & wait
 - No preemption
- ✦ Not always possible...

Eliminating mutual exclusion


- ✦ Some devices (such as printer) can be spooled
 - Only the printer daemon uses printer resource
 - This eliminates deadlock for printer
- ✦ Not all devices can be spooled
- ✦ Principle:
 - Avoid assigning resource when not absolutely necessary
 - As few processes as possible actually claim the resource

Attacking “hold and wait”

- ✦ Require processes to request resources before starting
 - A process never has to wait for what it needs
- ✦ This can present problems
 - A process may not know required resources at start of run
 - This also ties up resources other processes could be using
 - Processes will tend to be conservative and request resources they might need
- ✦ Variation: a process must give up all resources before making a new request
 - Process is then granted all prior resources as well as the new ones
 - Problem: what if someone grabs the resources in the meantime—how can the process save its state?


Attacking “no preemption”

- ✦ This is not usually a viable option
- ✦ Consider a process given the printer
 - Halfway through its job, take away the printer
 - Confusion ensues!
- ✦ May work for some resources
 - Forcibly take away memory pages, suspending the process
 - Process may be able to resume with no ill effects



Attacking “circular wait”

- ◆ Assign an order to resources
- ◆ Always acquire resources in numerical order
 - Need not acquire them all at once!
- ◆ Circular wait is prevented
 - A process holding resource n can't wait for resource m if $m < n$
 - No way to complete a cycle
 - Place processes above the highest resource they hold and below any they're requesting
 - All arrows point up!



Deadlock prevention: summary

- ✦ Mutual exclusion
 - Spool everything
- ✦ Hold and wait
 - Request all resources initially
- ✦ No preemption
 - Take resources away
- ✦ Circular wait
 - Order resources numerically

Example: two-phase locking

✦ Phase One

- Process tries to lock all data it needs, one at a time
- If needed data found locked, start over
- (no real work done in phase one)

✦ Phase Two

- Perform updates
- Release locks

✦ Note similarity to requesting all resources at once

✦ This is often used in databases

✦ It avoids deadlock by eliminating the “hold-and-wait” deadlock condition

“Non-resource” deadlocks

- ✦ Possible for two processes to deadlock
 - Each is waiting for the other to do some task
- ✦ Can happen with semaphores
 - Each process required to do a down() on two semaphores (mutex and another)
 - If done in wrong order, deadlock results
- ✦ Semaphores could be thought of as resources...

Starvation

- ✦ Algorithm to allocate a resource
 - Give the resource to the shortest job first
- ✦ Works great for multiple short jobs in a system
- ✦ May cause long jobs to be postponed indefinitely
 - Even though not blocked
- ✦ Solution
 - First-come, first-serve policy
- ✦ Starvation can lead to deadlock
 - Process starved for resources can be holding resources
 - If those resources aren't used and released in a timely fashion, shortage could lead to deadlock

Livelock

- ✦ Sometimes, processes can still run, but not make progress
- ✦ Example: two processes want to use resources A and B
 - P0 gets A, P1 gets B
 - Each realizes that a deadlock will occur if they proceed as planned!
 - P0 drops A, P1 drops B
 - P0 gets B, P1 gets A
 - Same problem as before
 - This can go on for a very long time...
- ✦ Real-world example: Ethernet transmission collisions
 - If there's a "collision" on the wire, wait and try again
 - Multiple processes waited the exact same amount of time...