

Part 5: I/O Systems

Input/Output

- ◆ Principles of I/O hardware
- ◆ Principles of I/O software
- ◆ I/O software layers
- ◆ Disks
- ◆ Clocks
- ◆ Character-oriented terminals
- ◆ Graphical user interfaces
- ◆ Network terminals
- ◆ Power management


How fast is I/O hardware?

Device	Data rate
Keyboard	10 bytes/sec
Mouse	100 bytes/sec
56K modem	7 KB/sec
Printer / scanner	200 KB/sec
USB	1.5 MB/sec
Digital camcorder	4 MB/sec
Fast Ethernet	12.5 MB/sec
Hard drive	20 MB/sec
FireWire (IEEE 1394)	50 MB/sec
XGA monitor	60 MB/sec
PCI bus	500 MB/sec


Device controllers

- ✦ I/O devices have components
 - Mechanical component
 - Electronic component
- ✦ Electronic component controls the device
 - May be able to handle multiple devices
 - May be more than one controller per mechanical component (example: hard drive)
- ✦ Controller's tasks
 - Convert serial bit stream to block of bytes
 - Perform error correction as necessary
 - Make available to main memory


Memory-Mapped I/O



Separate
I/O & memory
space



Memory-mapped I/O




Hybrid: both
memory-mapped &
separate spaces

How is memory-mapped I/O done?


◆ Single-bus

- All memory accesses go over a shared bus
- I/O and RAM accesses compete for bandwidth



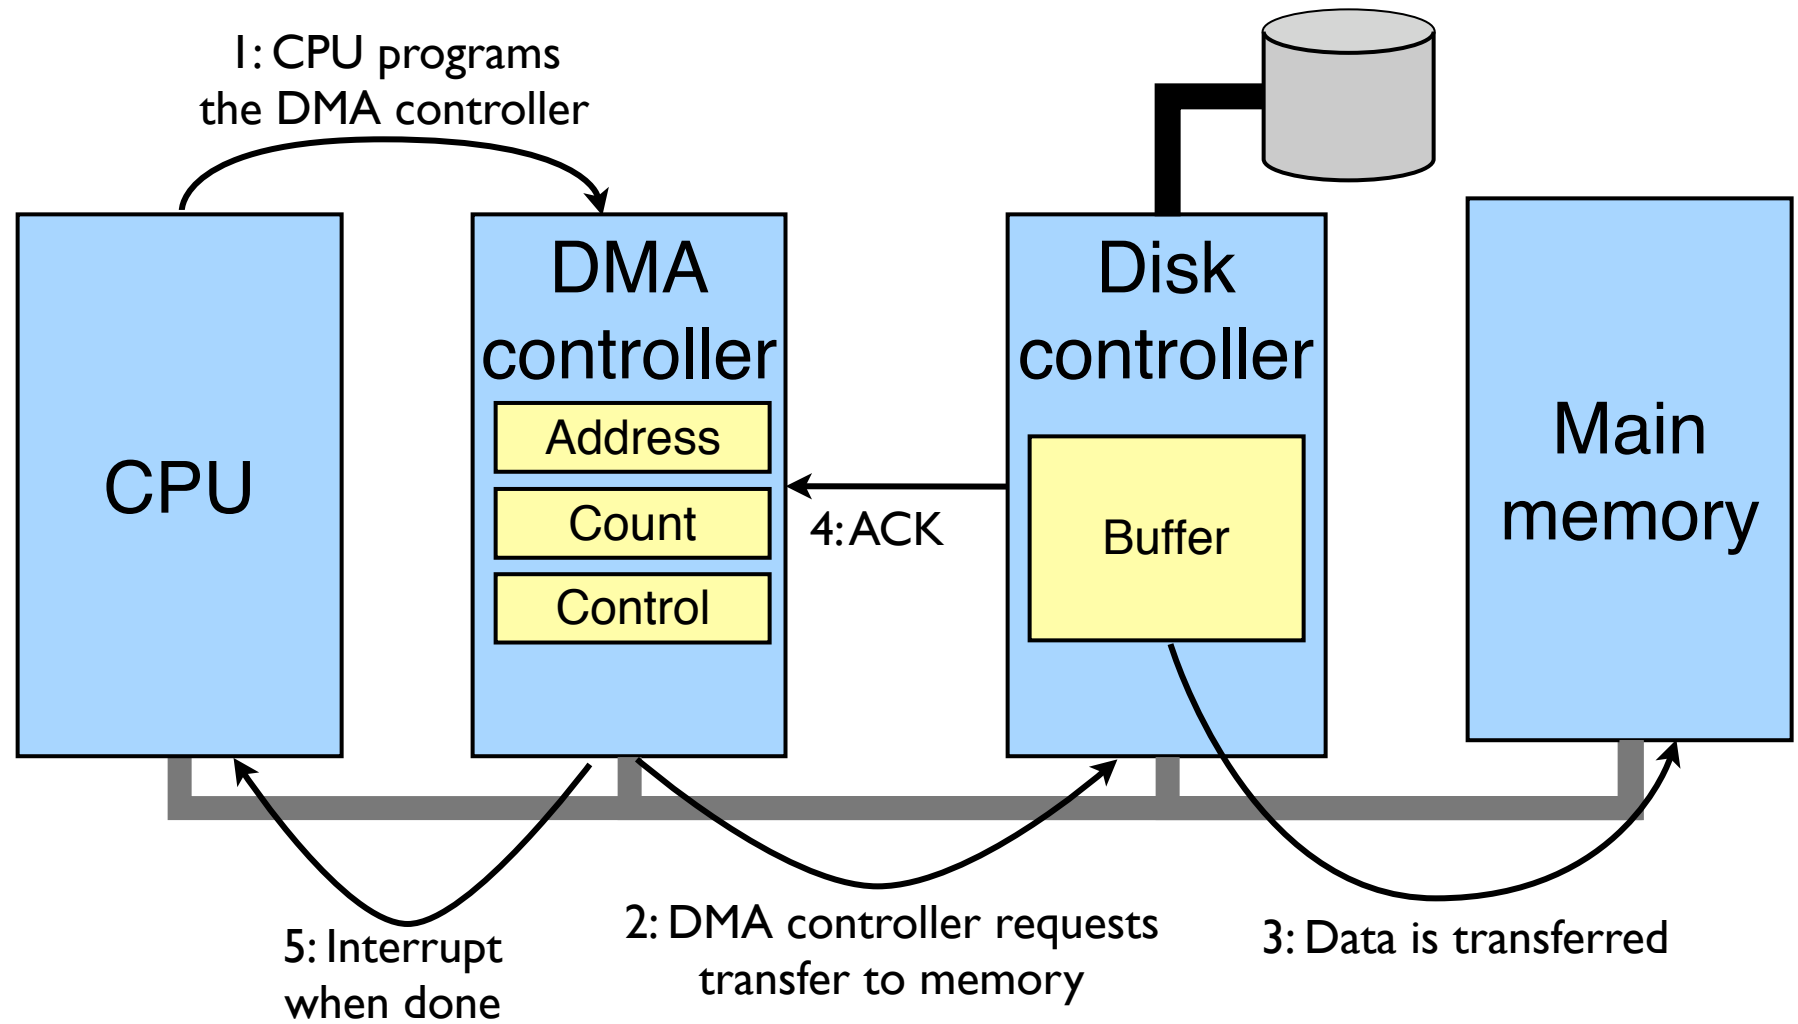
◆ Dual-bus

- RAM access over high-speed bus
- I/O access over lower-speed bus
- Less competition
- More hardware (more expensive...)

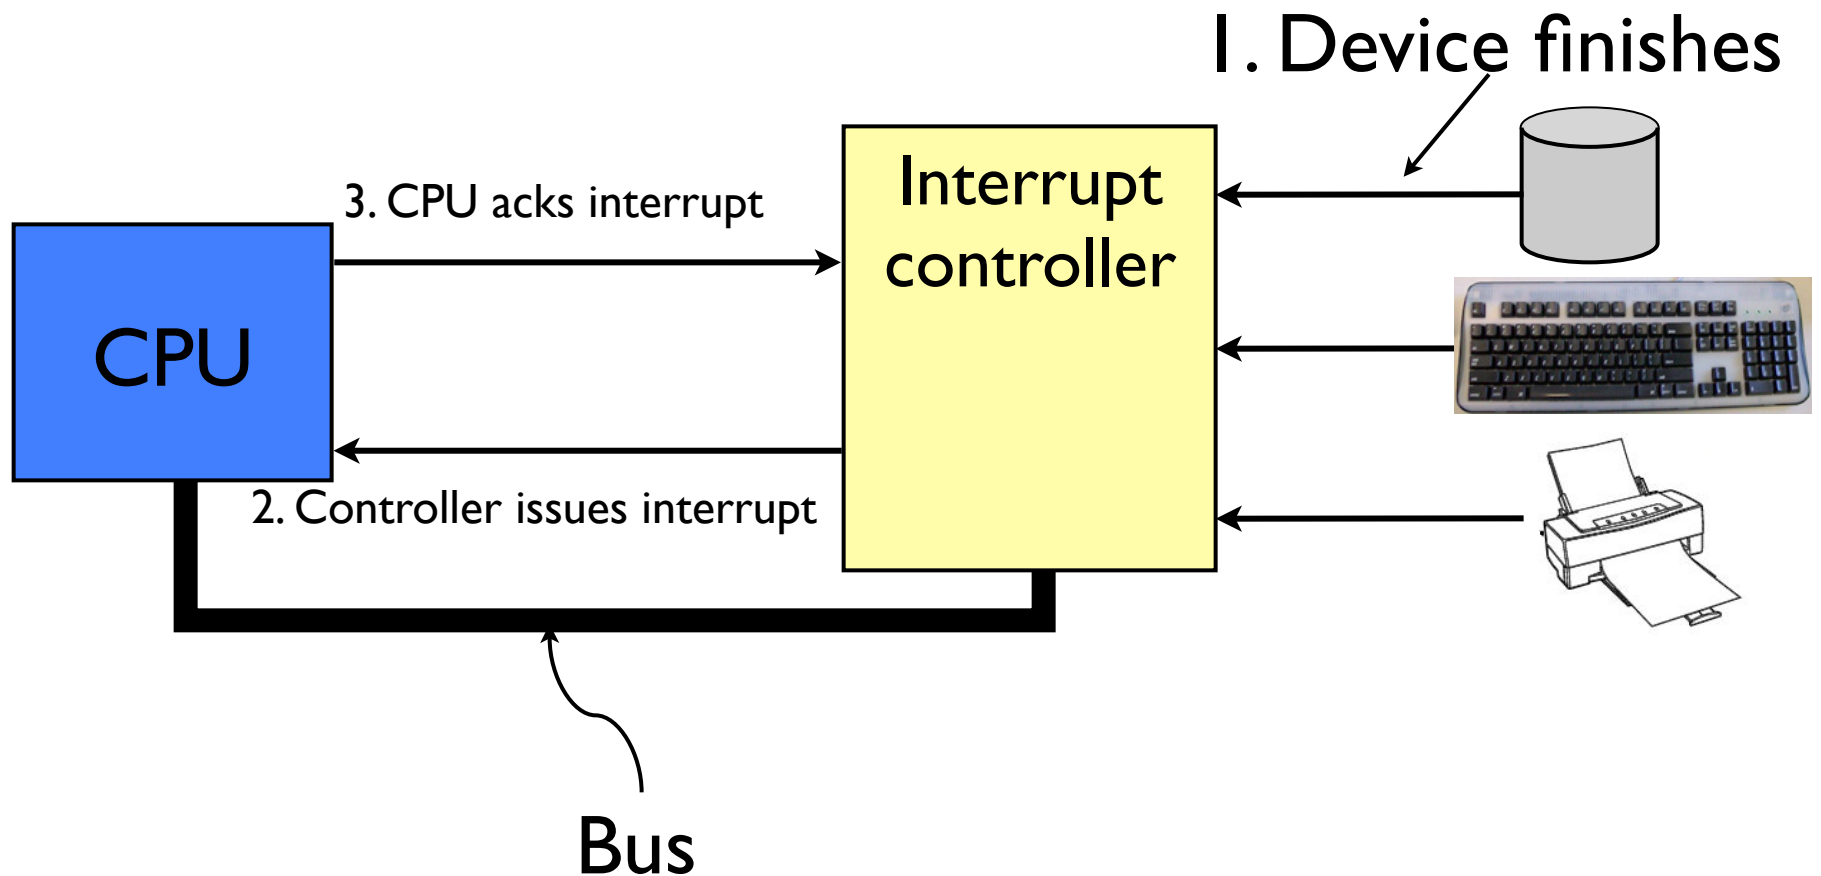


↖ This port allows I/O devices access into memory

Direct Memory Access (DMA)



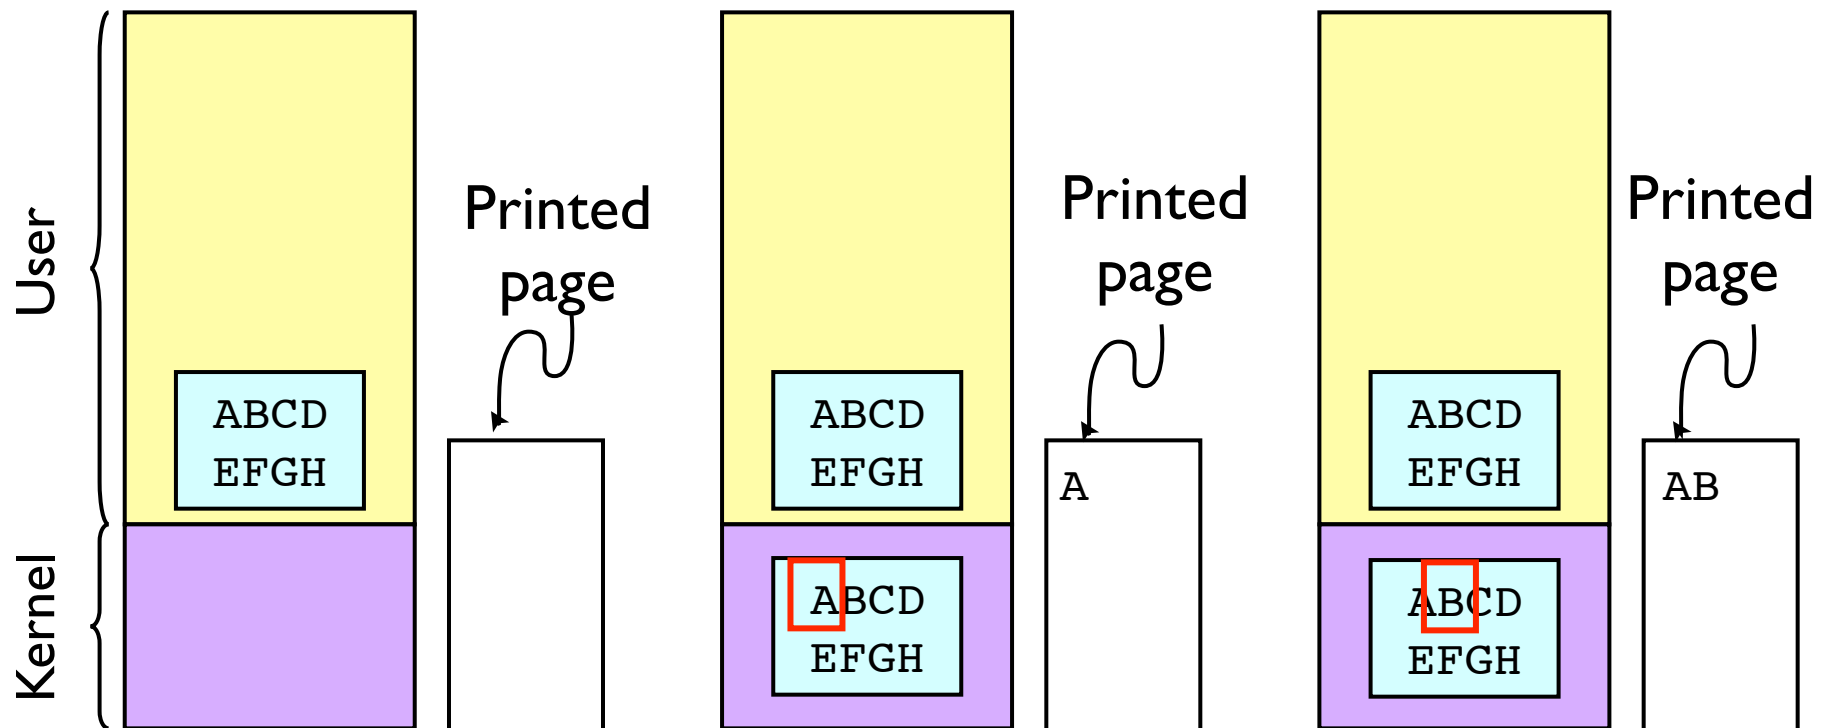
Hardware's view of interrupts



I/O software: goals

- ◆ Device independence
 - Programs can access any I/O device
 - No need to specify device in advance
- ◆ Uniform naming
 - Name of a file or device is a string or an integer
 - Doesn't depend on the machine (underlying hardware)
- ◆ Error handling
 - Done as close to the hardware as possible
 - Isolate higher-level software
- ◆ Synchronous vs. asynchronous transfers
 - Blocked transfers vs. interrupt-driven
- ◆ Buffering
 - Data coming off a device cannot be stored in final destination
- ◆ Sharable vs. dedicated devices

Programmed I/O: printing a page



Code for programmed I/O

```
copy_from_user (buffer, p, count); // copy into kernel buffer
for (j = 0; j < count; j++) {      // loop for each char
    while (*printer_status_reg != READY)
        ;                          // wait for printer to be ready
    *printer_data_reg = p[j];       // output a single character
}
return_to_user();
```

Interrupt-driven I/O

```
copy_from_user (buffer, p, count);  
j = 0;  
enable_interrupts();  
while (*printer_status_reg != READY)  
    ;  
*printer_data_reg = p[0];  
scheduler(); // and block user
```

```
if (count == 0) {  
    unblock_user();  
} else {  
    *printer_data_reg = p[j];  
    count--;  
    j++;  
}  
acknowledge_interrupt();  
return_from_interrupt();
```

Code run by
system call

Code run at
interrupt time

I/O using DMA

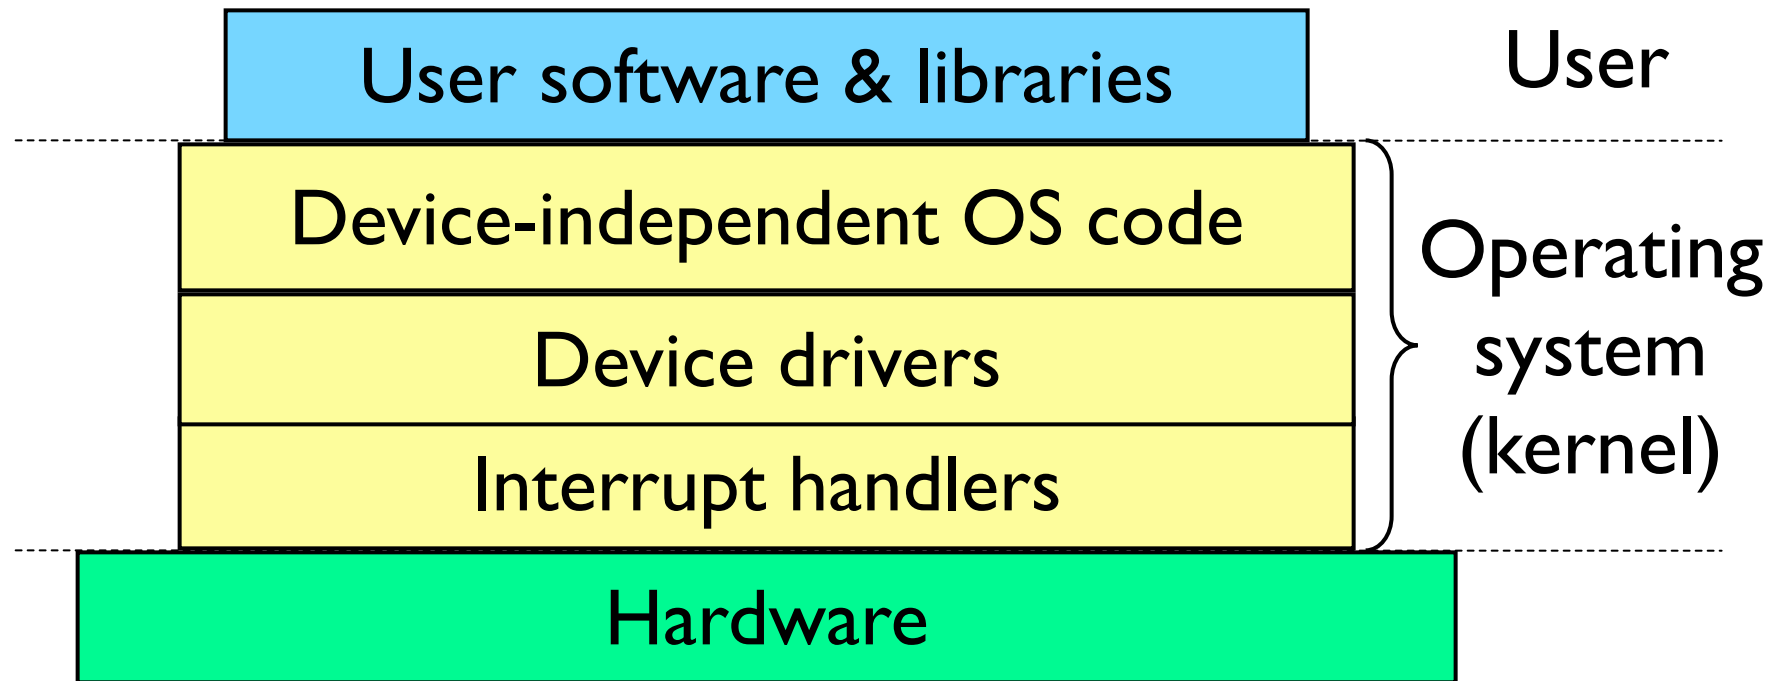
```
copy_from_user (buffer, p, count);  
set_up_DMA_controller();  
scheduler(); // and block user
```

Code run by
system call

```
acknowledge_interrupt();  
unlock_user();  
return_from_interrupt();
```

Code run at
interrupt time

Layers of I/O software



Interrupt handlers

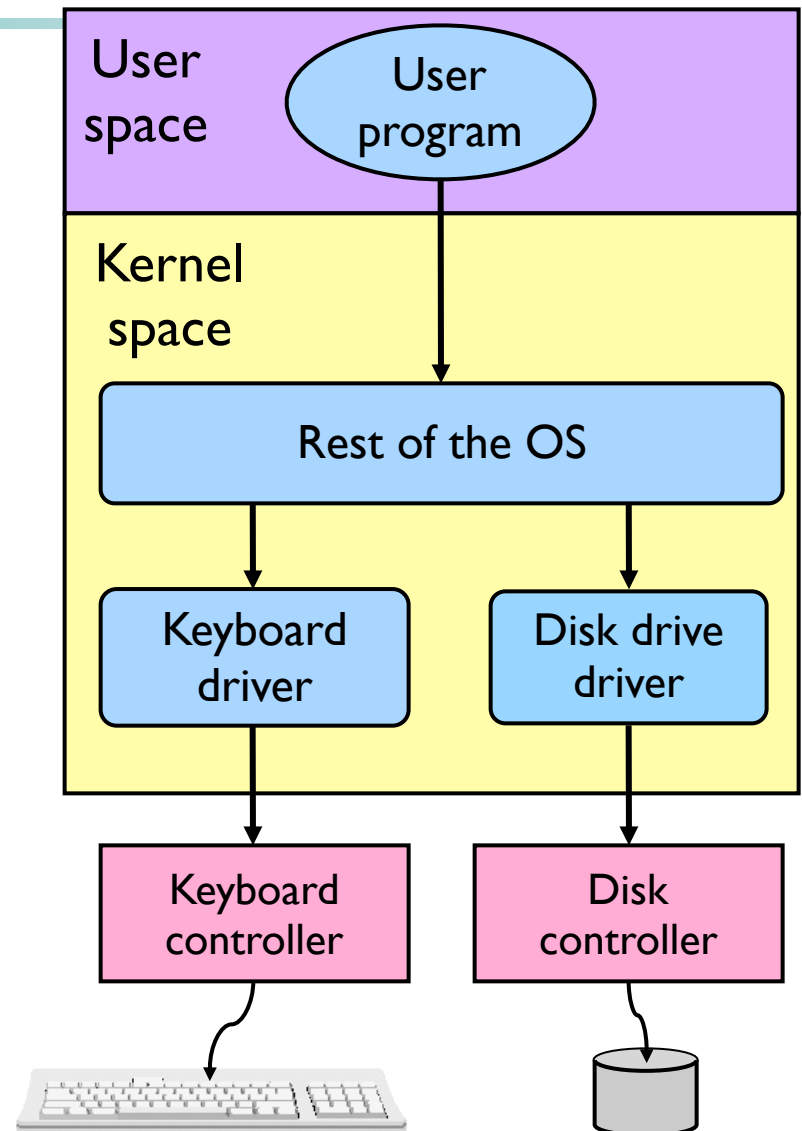
- ✦ Interrupt handlers are best hidden
 - Driver starts an I/O operation and blocks
 - Interrupt notifies of completion
- ✦ Interrupt procedure does its task
 - Then unblocks driver that started it
 - Perform minimal actions at interrupt time
 - Some of the functionality can be done by the driver after it is unblocked
- ✦ Interrupt handler must
 - Save registers not already saved by interrupt hardware
 - Set up context for interrupt service procedure

What happens on an interrupt

- ✦ Set up stack for interrupt service procedure
- ✦ Ack interrupt controller, reenale interrupts
- ✦ Copy registers from where saved
- ✦ Run service procedure
- ✦ (optional) Pick a new process to run next
- ✦ Set up MMU context for process to run next
- ✦ Load new process' registers
- ✦ Start running the new process

Device drivers


- ✦ Device drivers go between device controllers and rest of OS
 - Drivers standardize interface to widely varied devices
- ✦ Device drivers communicate with controllers over bus
 - Controllers communicate with devices themselves



Device-independent I/O software


- ✦ Device-independent I/O software provides common “library” routines for I/O software
- ✦ Helps drivers maintain a standard appearance to the rest of the OS
- ✦ Uniform interface for many device drivers for
 - Buffering
 - Error reporting
 - Allocating and releasing dedicated devices
 - Suspending and resuming processes
- ✦ Common resource pool
 - Device-independent block size (keep track of blocks)
 - Other device driver resources

Why a standard driver interface?



Non-standard device driver interface


- Different interface for each driver
- High operating system complexity
- Less code reuse



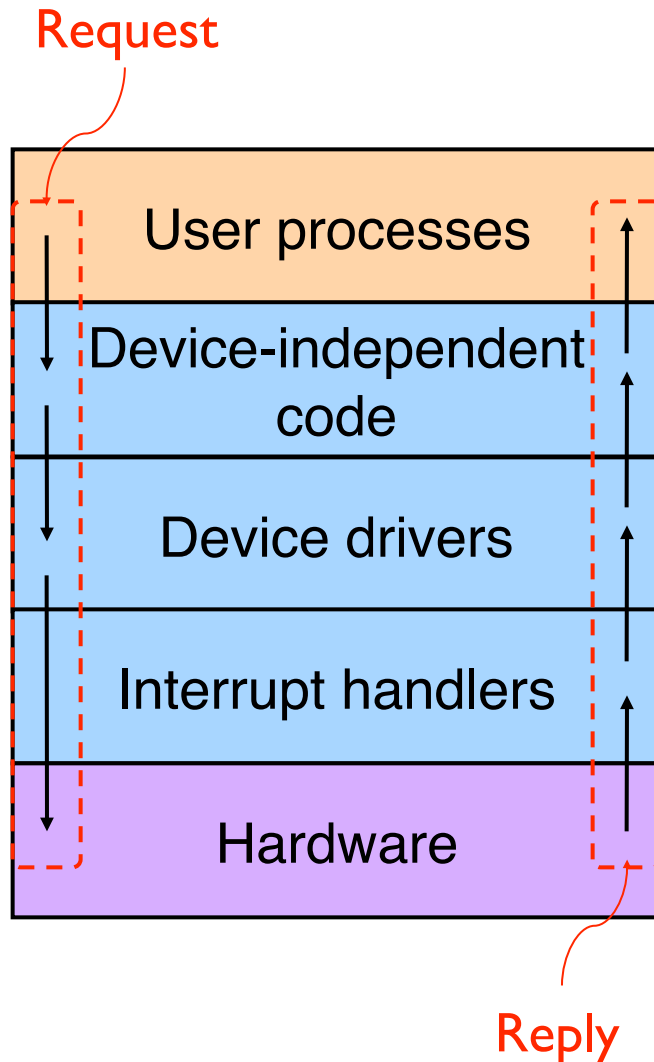
Standard device driver interface

- Less OS/driver interface code
- Lower OS complexity
- Easy to add new

Buffering device input



What happens where on an I/O request?



Make I/O call; format I/O; spooling

Naming, protection
blocking / buffering / allocation


Manage device registers & status

Signal device driver on completed I/O

Actually do the I/O (in hardware)

Disk drive structure

- ◆ Data stored on surfaces
 - Up to two surfaces per platter
 - One or more platters per disk
- ◆ Data in concentric tracks
 - Tracks broken into sectors
 - 256B-1KB per sector
 - Cylinder: corresponding tracks on all surfaces
- ◆ Data read and written by heads
 - Actuator moves heads
 - Heads move in unison




Disk drive specifics

	IBM 360KB floppy	Hitachi 500 GB HD
Cylinders	40	132,000 (estimate)
Tracks per cylinder	2	10
Sectors per track	9	750 (average)
Sectors per disk	720	10^9
Bytes per sector	512	512
Capacity	360 KB	500 GB
Seek time (minimum)	6 ms	2 ms (?)
Seek time (average)	77 ms	8.5 ms
Rotation time	200 ms	8.33 ms
Spinup time	250 ms	~5 sec (?)
Sector transfer time	22 ms	11 μ sec

Disk “zones”

- ✦ Outside tracks are longer than inside tracks
- ✦ Two options for longer tracks
 - Bits are “bigger”
 - More bits (transfer faster)
- ✦ Modern hard drives use second option
 - More data on outer tracks
- ✦ Disk divided into “zones”
 - Constant sectors per track in each zone
 - 8–30 (or more) zones on a disk




Disk “addressing”

- ✦ Millions of sectors on the disk must be labeled
- ✦ Two possibilities
 - Cylinder/track/sector
 - Sequential numbering
- ✦ Modern drives use sequential numbers
 - Disks map sequential numbers into specific location
 - Mapping may be modified by the disk
 - Remap bad sectors
 - Optimize performance
 - Hide the exact geometry, making life simpler for the OS


Building a better “disk”

- ✦ Problem: CPU performance has been increasing exponentially, but disk performance hasn't
 - Disks are limited by mechanics
- ✦ Problem: disks aren't all that reliable
- ✦ Solution: distribute data across disks, and use some of the space to improve reliability
 - Data transferred in parallel
 - Data stored across drives (striping)
 - Parity on an “extra” drive for reliability

RAIDs, RAIDs, and more RAIDs




CD-ROM recording



- ✦ CD-ROM has data in a spiral
 - Hard drives have concentric circles of data
- ✦ One continuous track: just like vinyl records!
- ✦ Pits & lands “simulated” with heat-sensitive material on CD-Rs and CD-RWs


Structure of a disk sector



- ✦ Preamble contains information about the sector
 - Sector number & location information
- ✦ Data is usually 256, 512, or 1024 bytes
- ✦ ECC (Error Correcting Code) is used to detect & correct minor errors in the data


Sector layout on disk

- ✦ Sectors numbered sequentially on each track
- ✦ Numbering starts in different place on each track: **sector skew**
 - Allows time for switching head from track to track
- ✦ All done to minimize delay in sequential transfers
- ✦ In modern drives, this is only approximate!




Sector interleaving


- ✦ On older systems, the CPU was slow => time elapsed between reading consecutive sectors
- ✦ Solution: leave space between consecutively numbered sectors
- ✦ This isn't done much these days...



No interleaving



Skipping 1 sector



Skipping 2 sectors

What's in a disk request?

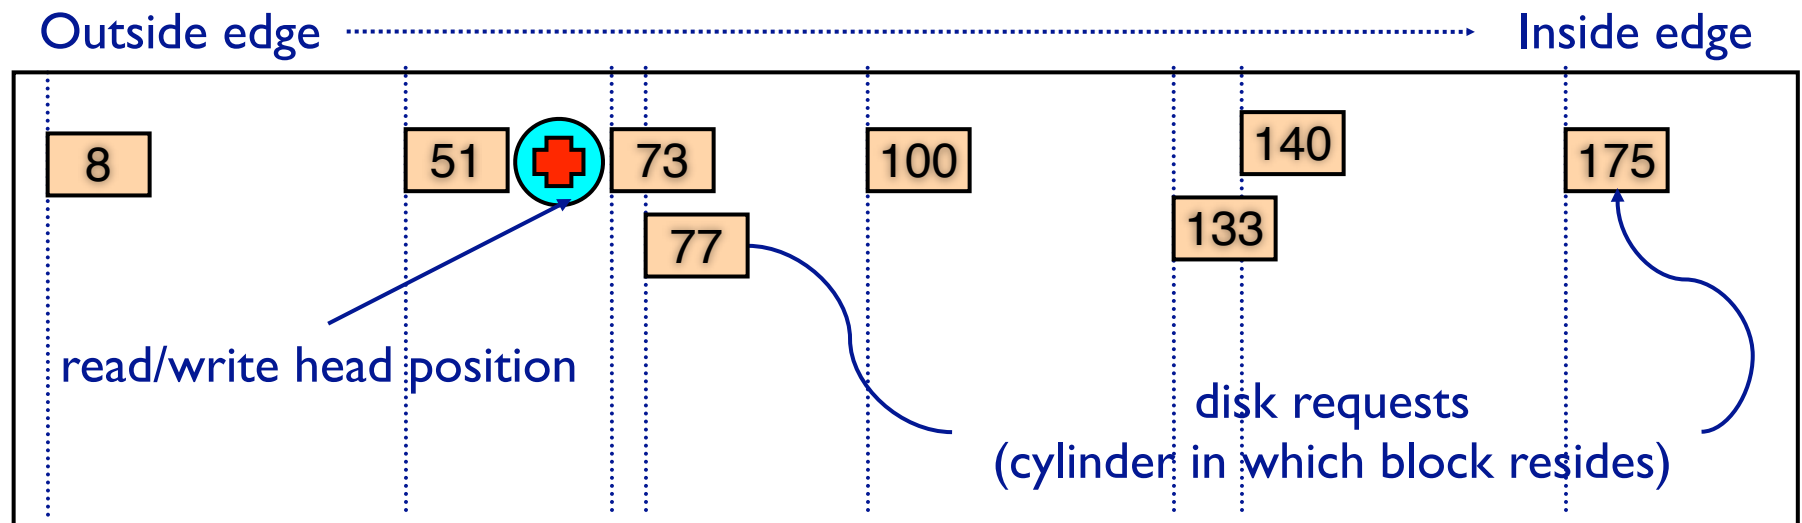
- ✦ Time required to read or write a disk block determined by 3 factors
 - Seek time
 - Rotational delay
 - Average delay = $1/2$ rotation time
 - Example: rotate in 10 ms, average rotation delay = 5 ms
 - Actual transfer time
 - Transfer time = time to rotate over sector
 - Example: rotate in 10 ms, 200 sectors/track $\Rightarrow 10/200$
ms = 0.05 ms transfer time per sector
- ✦ Seek time dominates, with rotation time close
- ✦ Error checking is done by controllers

Disk request scheduling

- ◆ Goal: use disk hardware efficiently
 - Bandwidth as high as possible
 - Disk transferring as often as possible (and not seeking)
- ◆ We want to
 - Minimize disk seek time (moving from track to track)
 - Minimize rotational latency (waiting for disk to rotate the desired sector under the read/write head)
- ◆ Calculate disk bandwidth by
 - Total bytes transferred / time to service request
 - Seek time & rotational latency are overhead (no data is transferred), and reduce disk bandwidth
- ◆ Minimize seek time & rotational latency by
 - Using algorithms to find a good sequence for servicing requests
 - Placing blocks of a given file “near” each other

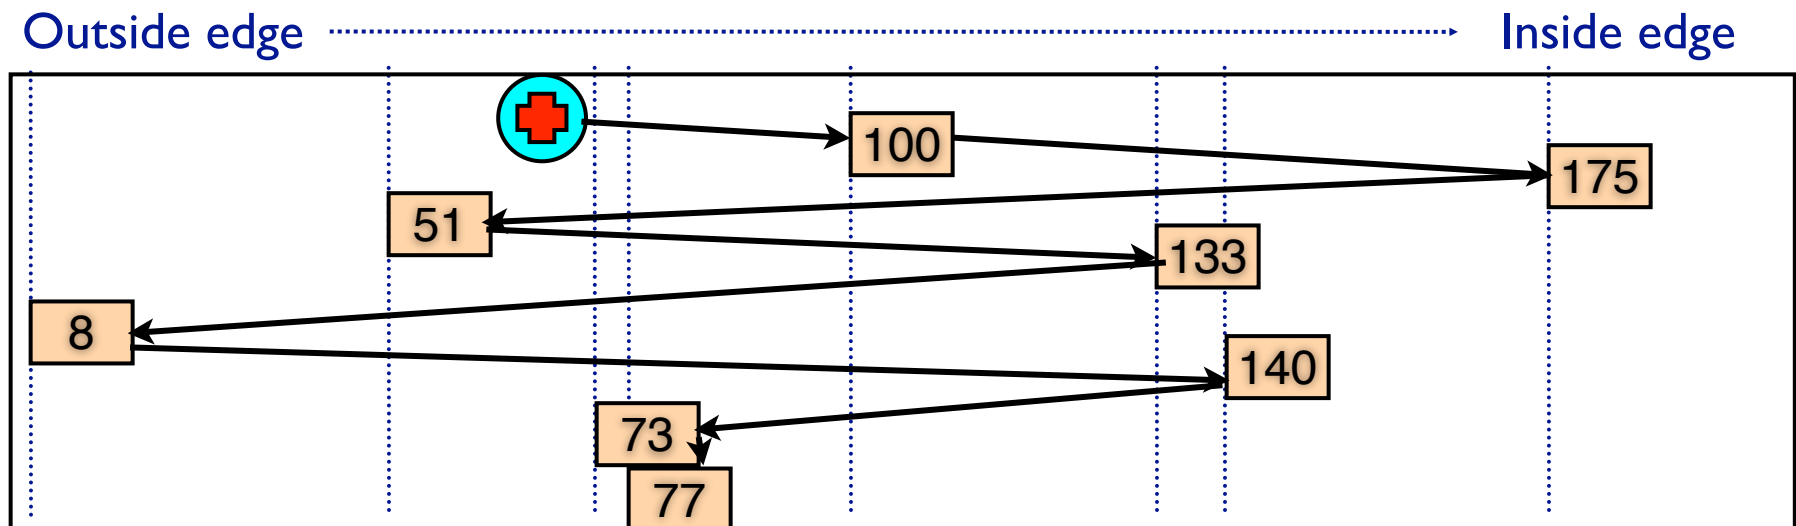
Disk scheduling algorithms

- ◆ Schedule disk requests to minimize disk seek time
 - Seek time increases as distance increases (though not linearly)
 - Minimize seek distance -> minimize seek time
- ◆ Disk seek algorithm examples assume a request queue & head position (disk has 200 cylinders)
 - Queue = 100, 175, 51, 133, 8, 140, 73, 77
 - Head position = 63



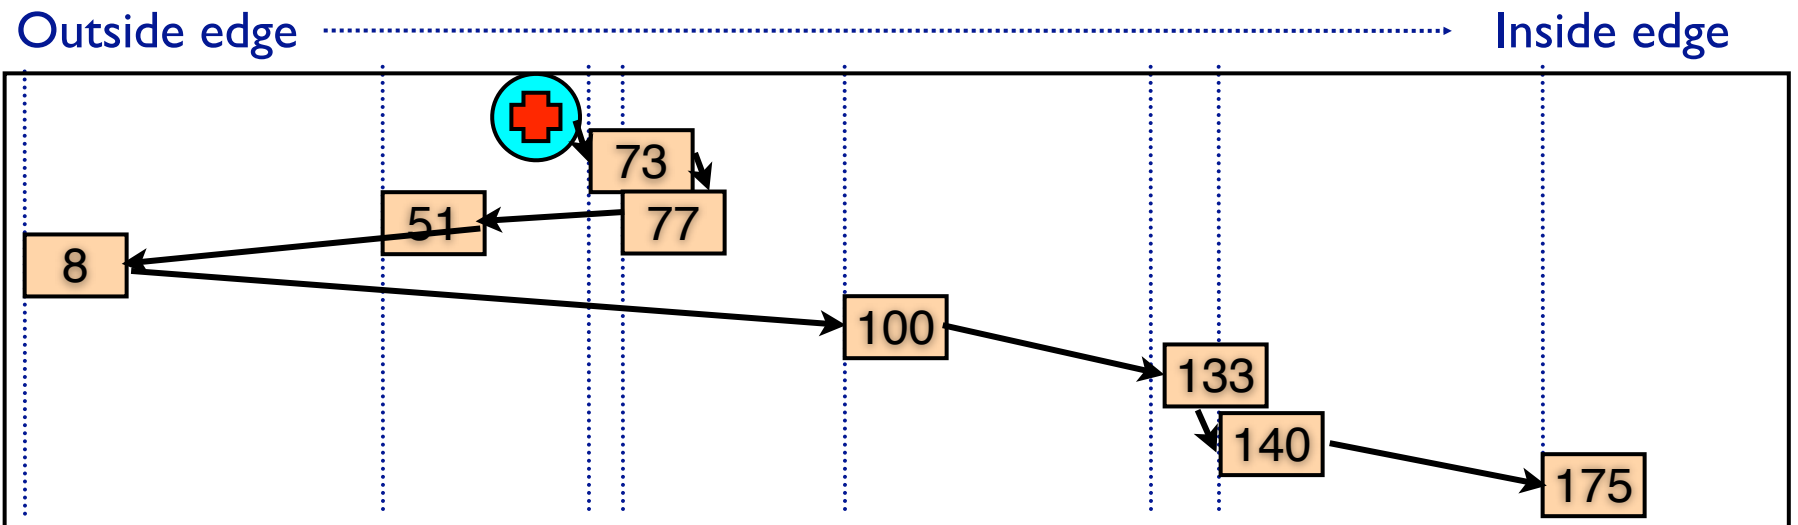
First-Come-First Served (FCFS)

- ◆ Requests serviced in the order in which they arrived
 - Easy to implement!
 - May involve lots of unnecessary seek distance
- ◆ Seek order = 100, 175, 51, 133, 8, 140, 73, 77
- ◆ Seek distance = $(100-63) + (175-100) + (175-51) + (133-51) + (133-8) + (140-8) + (140-73) + (77-73) = 646$ cylinders



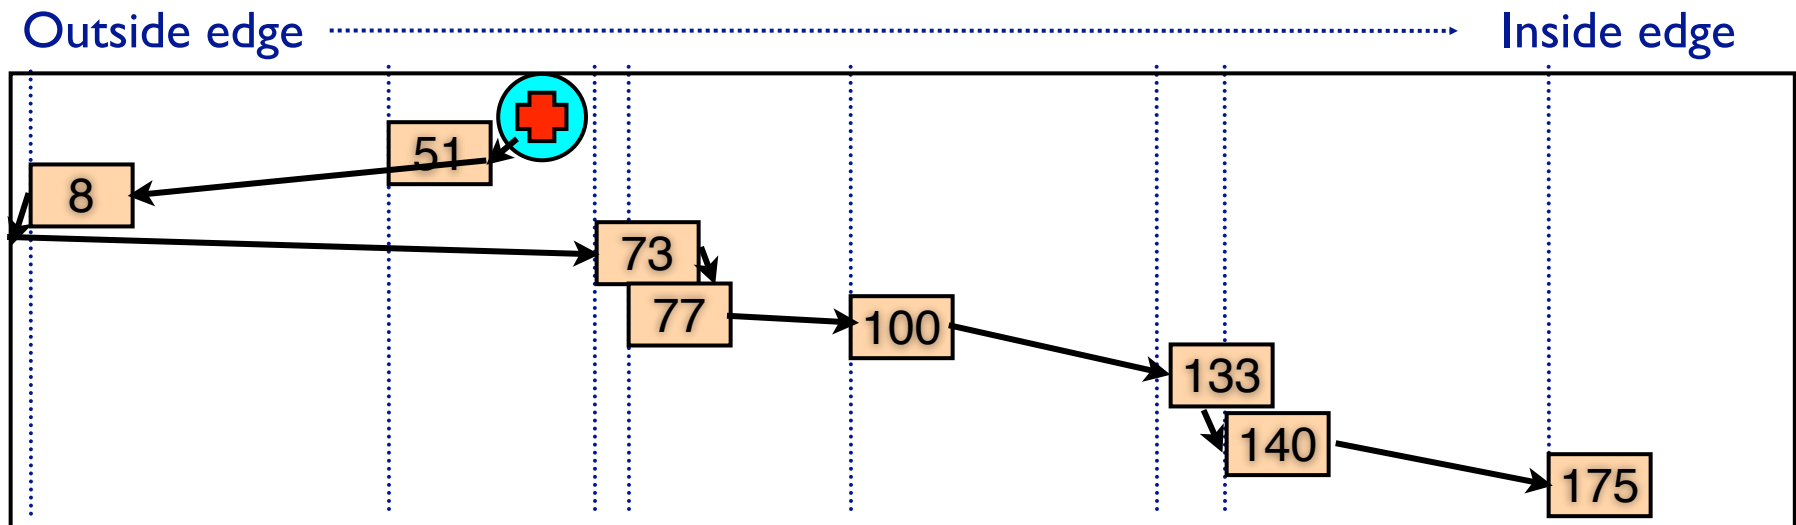
Shortest Seek Time First (SSTF)

- ◆ Service the request with the shortest seek time from the current head position
 - Form of SJF scheduling
 - May starve some requests
- ◆ Seek order = 73, 77, 51, 8, 100, 133, 140, 175
- ◆ Seek distance = $10 + 4 + 26 + 43 + 92 + 33 + 7 + 35 = 250$ cylinders



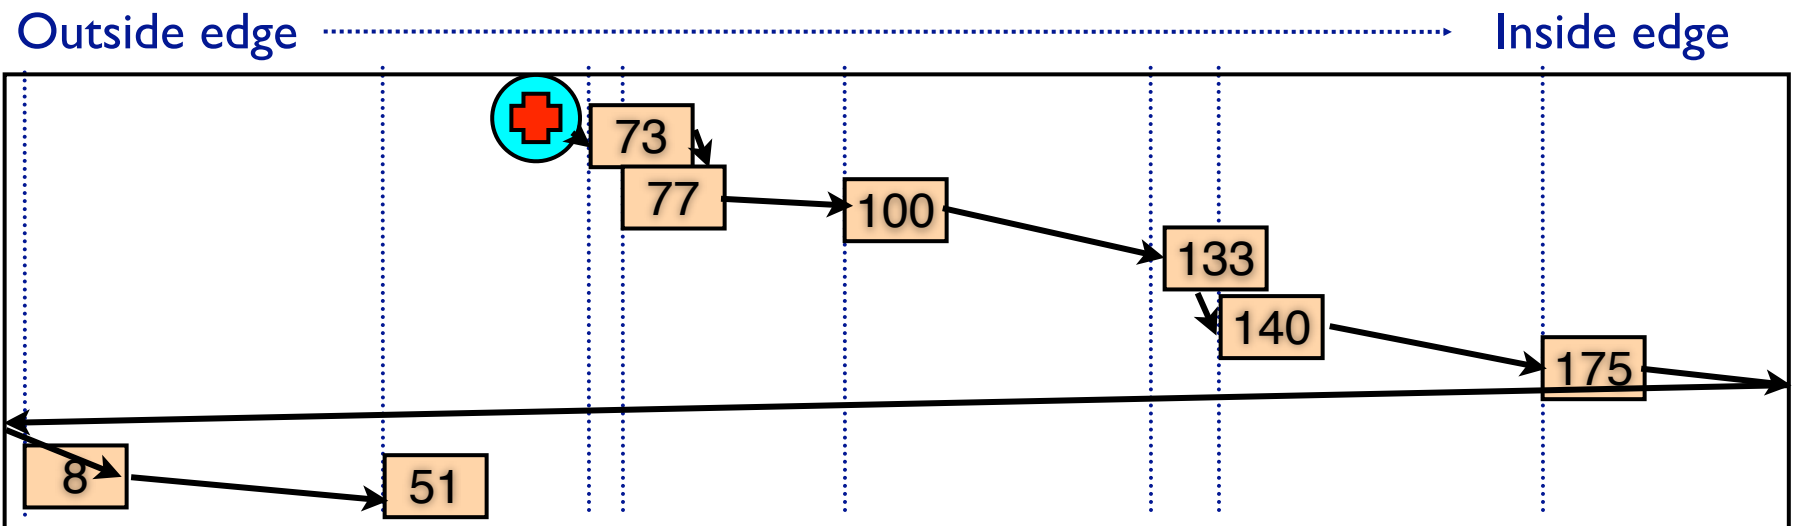
SCAN (elevator algorithm)

- ◆ Disk arm starts at one end of the disk and moves towards the other end, servicing requests as it goes
 - Reverses direction when it gets to end of the disk
 - Also known as elevator algorithm
- ◆ Seek order = 51, 8, 0, 73, 77, 100, 133, 140, 175
- ◆ Seek distance = $12 + 43 + 8 + 73 + 4 + 23 + 33 + 7 + 35 = 238$ cyls



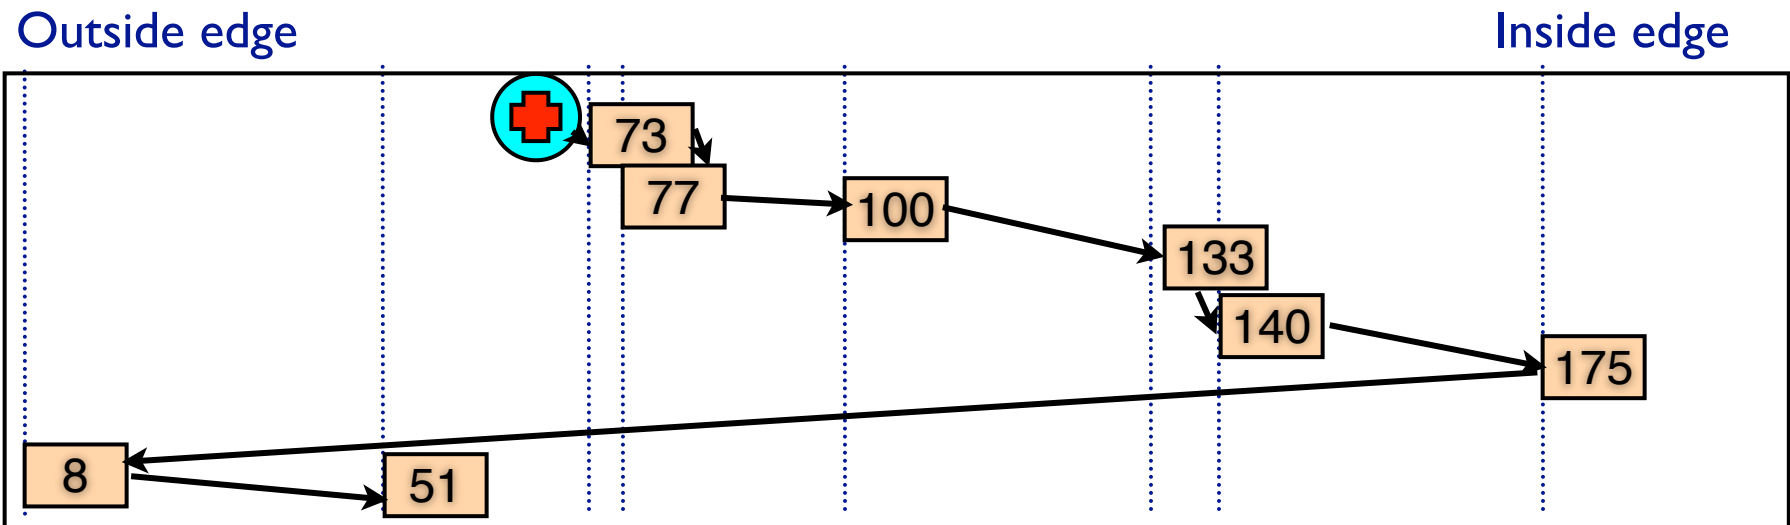
C-SCAN

- ◆ Identical to SCAN, except head returns to cylinder 0 when it reaches the end of the disk
 - Treats cylinder list as a circular list that wraps around the disk
 - Waiting time is more uniform for cylinders near the edge of the disk
- ◆ Seek order = 73, 77, 100, 133, 140, 175, 199, 0, 8, 51
- ◆ Distance = 10 + 4 + 23 + 33 + 7 + 35 + 24 + 199 + 8 + 43



C-LOOK

- ✦ Identical to C-SCAN, except head only travels as far as the last request in each direction
 - Saves seek time from last sector to end of disk
- ✦ Seek order = 73, 77, 100, 133, 140, 175, 8, 51
- ✦ Distance = $10 + 4 + 23 + 33 + 7 + 35 + 167 + 43 = 322$ cylinders




Picking a disk scheduling algorithm


- ✦ SSTF is easy to implement and works OK if there aren't too many disk requests in the queue
- ✦ SCAN-type algorithms perform better for systems under heavy load
 - More fair than SSTF
 - Use LOOK rather than SCAN algorithms to save time
- ✦ Long seeks aren't too expensive, so choose C-LOOK over LOOK to make response time more even
- ✦ Disk request scheduling interacts with algorithms for allocating blocks to files
 - Make scheduling algorithm modular: allow it to be changed without changing the file system
- ✦ Use SSTF for lightly loaded systems
- ✦ Use C-LOOK for heavily loaded systems

When good disks go bad...


- ◆ Disks have defects
 - In nearly a billion sectors, this isn't surprising!
- ◆ ECC helps with errors, but sometimes this isn't enough
- ◆ Disks keep spare sectors (normally unused) and remap bad sectors into these spares
 - If there's time, the whole track could be reordered...



(a)



(b)



(c)

Flash memory

- ✦ Compared to disk, flash is
 - Faster (shorter access time, but lower bandwidth)
 - More expensive
- ✦ Compared to DRAM, flash is
 - Cheaper (a bit)
 - Non-volatile (data survives power loss)
 - Slower
- ✦ Use flash as a level between disk and memory?
- ✦ Problems
 - Flash wears out: can only write 10,000–100,000 times per memory cell
 - Flash isn't too reliable: lots of bit errors

Handling flash in the OS

- ✦ Treat it like a disk?
 - Flash is written in blocks, just like a disk
 - Blocks have to be erased first: somewhat slow
 - Flash is random access, just like a disk
 - ➡ This approach is often used!
- ✦ Need to be careful about wearing out flash
 - Most flash devices do “wear leveling”: remap blocks internally when they’re erased
 - File systems for flash should take wear into account
 - Many don’t, including the standard VFAT file system

Clock hardware

- ✦ Crystal oscillator with fixed frequency (only when computer is on)
 - Typically used to time short intervals (~ 1 second)
 - May be used to correct time-of-day clock
- ✦ Time-of-day clock (runs when system is off)
 - Keeps minutes, hours, days
 - May not be too accurate...
 - Used to load system clock at startup
- ✦ Time kept in seconds and ticks (often 100–1000 per second)
 - Number of seconds since a particular time
 - For many versions of Unix, tick 0 was on January 1, 1970
 - Number of ticks this second
 - Advance ticks once per clock interrupt
 - Advance seconds when ticks “overflow”


Keeping time

- ✦ Check time via the Web
 - Protocol to get time from accurate time servers
- ✦ What happens when system clock is slow?
 - Advance clock to the correct current time
 - May be done all at once or over a minute or two
- ✦ What happens when clock is fast?
 - Can't have time run backwards!
 - Instead, advance time more slowly than normal until the clock is correct
- ✦ Track clock drift, do periodic updates to keep clock accurate

Using timers in software


- ◆ Use short interval clock for timer interrupts
 - Specified by applications
 - No problems if interrupt frequency is low
 - May have multiple timers using a single clock chip
- ◆ Use soft timers to avoid interrupts
 - Kernel checks for soft timer expiration before it exits to user mode
 - Less accurate than using a hardware timer
 - How well this works depends on rate of kernel entries

Current time: 5100





Where does the power go?

- ◆ How much power does each part of a laptop computer use?
 - Two studies: 1994 & 1998
 - Most power to the display!
- ◆ Save power by
 - Reducing display power
 - Slowing down CPU
 - Cutting power used by disk




1994



1998

Reducing CPU power usage




- ◆ Running at full clock speed
 - Jobs finish quickly
 - High energy consumption: proportional to shaded area
 - Intel processors may use 50–75 watts!
- ◆ Cutting voltage by two
 - Cuts clock speed by two: processes take longer
 - Cuts power by four
 - Cuts energy consumption ($= \text{power} * \text{time}$) in half


How can we reduce power usage?

- ✦ Telling the programs to use less energy
 - May mean poorer user experience
 - Makes batteries last longer!
- ✦ Examples
 - Change from color output to black and white
 - Speech recognition reduces vocabulary
 - Less resolution or detail in an image
 - Fewer image updates per second

Holographic Storage



MEMS Storage



MEMS Storage

